

Distributed Systems Project Report

Course: DOS

Student: Ameed Jabr

Student#: 12043106

Project: Catalog Service, Order Service, and Frontend Service

Semester: first semester (fall) 2025

Github link: <https://github.com/Ameedjabr/Bazarcom>

1. Introduction

In this lab, we extend the system developed in Lab 1 by introducing **replication**, **load balancing**, and **in-memory caching** at the front-end server to improve system performance and scalability. The catalog and order services are replicated, while the front-end server remains a single entry point for all client requests.

The objectives of this lab are:

- Reduce request latency using caching
- Distribute load across replicated backend services
- Maintain cache consistency during write operations
- Measure and evaluate system performance with and without caching

2. System Architecture

2.1 Components

- **Front-End Service**
 - Receives all client requests
 - Implements round-robin load balancing
 - Contains an in-memory LRU cache for catalog read requests
- **Catalog Service (Replicated)**
 - Stores book information in a CSV file
 - Handles search and info requests

- Handles update requests (writes)
- **Order Service (Replicated)**
 - Handles purchase requests
 - Updates stock by communicating with catalog service

2.2 Replication

- Two catalog replicas (ports 5000, 5001)
- Two order replicas (ports 7000, 7001)
- Front-end distributes requests using **round-robin load balancing**

2.3 Caching

- Cache located in the **front-end server**
- Caches only **read requests** (/info/{id})
- Uses **LRU replacement policy**
- Cache entries are invalidated upon writes (purchase or update)

3. Caching Design

3.1 Cache Placement

The cache is integrated inside the front-end server as an in-memory data structure.

3.2 Cache Policy

- **Key:** Book ID
- **Value:** JSON response from catalog /info/{id}
- **Replacement Policy:** Least Recently Used (LRU)
- **Capacity:** 50 entries

3.3 Cache Consistency

To maintain strong consistency:

1. Write operations (purchase or catalog update) occur on the backend
2. The backend sends an **invalidate request** to the front-end
3. The corresponding cache entry is removed before future reads

4. Performance Experiments

Performance was measured using PowerShell's Measure-Command to issue multiple requests and compute average response times.

5. Experiment 1: Query and Buy Response Time

5.1 Methodology

- 30 repeated requests were sent for each operation
- Two scenarios were tested:
 - **With caching enabled**
 - **Without caching enabled**
- Endpoints tested:
 - Query: /info/5
 - Buy: /purchase/5

5.2 Results

Table 1: Query (Read) Response Time

Scenario	Average (ms)	Min (ms)	Max (ms)
Without Cache	12.93	9.05	95.18
With Cache	11.86	5.96	94.40

Table 2: Buy (Write) Response Time

Scenario	Average (ms)	Min (ms)	Max (ms)
Without Cache	14.66	8.55	128.23
With Cache	23.78	8.65	153.95

5.3 Analysis

- **Query requests benefit from caching**, as repeated reads are served directly from memory instead of contacting the catalog service.
- **Buy requests are slower with caching**, because:
 - Writes trigger cache invalidation
 - Additional communication occurs between services
- Caching improves **read latency**, while writes incur **consistency overhead**
 - **Answer (Q1):**

From the measurements, the average query response time was 12.93 ms without caching and 11.86 ms with caching, which is an improvement of 1.07 ms (~8.3% faster).
For buy requests, the average response time was 14.66 ms without caching and 23.78 ms with caching (slower), because buy operations trigger additional steps (catalog update + cache invalidation) to ensure strong consistency. Therefore, caching mainly helps read/query operations, while write/buy operations can become slower due to consistency overhead.
 - Caching improvement (query) = $(12.93 - 11.86) / 12.93 = \text{~8.3\%}$

6. Experiment 2: Cache Invalidation and Consistency Overhead

6.1 Methodology

1. A cached /info/5 request is issued (cache hit)
2. A /purchase/5 request is issued
3. The cache entry for item 5 is invalidated
4. A subsequent /info/5 request results in a **cache miss**

6.2 Observations

- Cache invalidation introduces extra network overhead
- The first read after a write experiences higher latency due to a cache miss

- Subsequent reads return to low latency once cached again

6.3 Results

Table 3: Cache Consistency Overhead

Operation	Approx. Latency (ms)
Cache hit (read)	~6–12 ms
Cache invalidation (write)	~18–25 ms
Cache miss after write	~40–70 ms

7. Load Balancing Verification

Round-robin load balancing was verified by:

- Running multiple replicas on different ports
- Observing correct responses even when only one replica is active
- Demonstrating correct request routing after enabling replicas

8. Conclusion

This lab demonstrates that:

- **Caching significantly improves read performance**
- **Replication and load balancing increase system scalability**
- **Strong cache consistency introduces overhead for write operations**
- A front-end cache is effective for read-heavy workloads

Overall, the system achieves improved performance while maintaining correctness and consistency.

Important screenshots

1. Front-end service running (terminal)

```
PS C:\Users\ME\Desktop\Bazarcom\frontend-service> mvn exec:java
[INFO] Scanning for projects...
[INFO]
[INFO] -----< bazar:frontend-service >-----
[INFO] Building frontend-service 1.0-SNAPSHOT
[INFO]   from pom.xml
[INFO] -----[ jar ]-----
[INFO]
[INFO] --- exec:3.1.0:java (default-cli) @ frontend-service ---
FrontEndService running at http://localhost:9000
```

2. Catalog service running (terminal)

```
[INFO] -----
○ PS C:\Users\ME\Desktop\Bazarcom\catalog-service> mvn exec:java
[INFO] Scanning for projects...
[INFO]
[INFO] -----< bazar:catalog-service >-----
[INFO] Building catalog-service 1.0-SNAPSHOT
[INFO]   from pom.xml
[INFO] -----[ jar ]-----
[INFO]
[INFO] --- exec:3.1.0:java (default-cli) @ catalog-service ---
Loaded 7 items from CSV
CatalogService running on http://localhost:5000
```

3. Order service running (terminal)

```
Warning: PowerShell detected that you might be using a screen reader and has disabled PSReadLine.
Run 'Import-Module PSReadLine'.
```

```
PS C:\Users\ME\Desktop\Bazarcom\order-service> mvn exec:java
[INFO] Scanning for projects...
[INFO]
[INFO] -----< bazar:order-service >-----
[INFO] Building order-service 1.0-SNAPSHOT
[INFO]   from pom.xml
[INFO] -----[ jar ]-----
[INFO]
[INFO] --- exec:3.1.0:java (default-cli) @ order-service ---
OrderService running at http://localhost:7000/order?id=1
```

Replication & Load Balancing

4. Catalog replica running on port 5000

```
[INFO] -----  
○ PS C:\Users\ME\Desktop\Bazarcom\catalog-service> mvn exec:java  
[INFO] Scanning for projects...  
[INFO]  
[INFO] -----< bazar:catalog-service >-----  
[INFO] Building catalog-service 1.0-SNAPSHOT  
[INFO]   from pom.xml  
[INFO] -----[ jar ]-----  
[INFO]  
[INFO] --- exec:3.1.0:java (default-cli) @ catalog-service ---  
Loaded 7 items from CSV  
CatalogService running on http://localhost:5000
```

5. Catalog replica running on port 5001

```
PS C:\Users\ME\Desktop\Bazarcom\catalog-service> mvn exec:java "-Dexec.args=5001"  
[INFO] Scanning for projects...  
[INFO]  
[INFO] -----< bazar:catalog-service >-----  
[INFO] Building catalog-service 1.0-SNAPSHOT  
[INFO]   from pom.xml  
[INFO] -----[ jar ]-----  
[INFO]  
[INFO] --- exec:3.1.0:java (default-cli) @ catalog-service ---  
Loaded 7 items from CSV  
CatalogService running on http://localhost:5001
```

6. Successful /search request through front-end

localhost:5001/search/undergraduate%20school

Pretty-print

```
{
  "items": [
    {
      "id": "3",
      "title": "Xen and the Art of Surviving Undergraduate School"
    },
    {
      "id": "4",
      "title": "Cooking for the Impatient Undergrad"
    },
    {
      "id": "5",
      "title": "How to finish Project 3 on time"
    },
    {
      "id": "6",
      "title": "Why theory classes are so hard"
    },
    {
      "id": "7",
      "title": "Spring in the Pioneer Valley"
    }
  ]
}
```


localhost:5000/search/undergraduate%20school

Pretty-print


```
{
  "items": [
    {
      "id": "3",
      "title": "Xen and the Art of Surviving Undergraduate School"
    },
    {
      "id": "4",
      "title": "Cooking for the Impatient Undergrad"
    },
    {
      "id": "5",
      "title": "How to finish Project 3 on time"
    },
    {
      "id": "6",
      "title": "Why theory classes are so hard"
    },
    {
      "id": "7",
      "title": "Spring in the Pioneer Valley"
    }
  ]
}
```

Caching

7. /info/5 request (first time – cache miss)



8. /info/5 request (second time – cache hit)



Cache Invalidation

9. /purchase/5 request success

```
{  
  "status": "success",  
  "message": "Book purchased successfully",  
  "title": "How to finish Project 3 on time",  
  "price": 30  
}
```

10. /info/5 request after purchase (cache miss)

Performance Measurement

11. PowerShell output – query with cache

```
Warning: PowerShell detected that you might be using a screen reader and has disabled PSReadLine for compatibility purposes. If you want to re-enable it, run 'Import-Module PSReadLine'.
```

```
PS C:\Users\ME\Desktop\Bazarcom> $times=@()
PS C:\Users\ME\Desktop\Bazarcom> 1..30 | % {
>>   $ms = (Measure-Command { Invoke-WebRequest -UseBasicParsing "http://localhost:9000/info/5" }).TotalMilliseconds
>>   $times += $ms
>> }
>> $times | Measure-Object -Average -Minimum -Maximum
>>
```

```
Count      : 30
Average    : 14.65655333333333
Sum        :
Maximum    : 128.2336
Minimum    : 8.5454
Property   :
```

```
PS C:\Users\ME\Desktop\Bazarcom> █
```

12. PowerShell output – query without cache

```
Count      : 30
Average    : 12.92590666666667
Sum        :
Maximum    : 95.1761
Minimum    : 9.051
Property   :
```

13. PowerShell output – buy with cache

```
Count      : 20
Average    : 23.78468
Sum        :
Maximum    : 153.9493
Minimum    : 8.6469
Property   :
```

14. PowerShell output – buy without cache

```
Count      : 20
Average    : 11.860335
Sum        :
Maximum    : 94.4014
Minimum    : 5.9603
Property   :
```